

Abstract

The mental contexts in which we interpret experiences are often person-specific, even when the experiences themselves are shared. We developed a geometric framework for mathematically characterizing the subjective conceptual content of dynamic naturalistic experiences. We model experiences and memories as “trajectories” through word embedding spaces whose coordinates reflect the universe of thoughts under consideration. Memory encoding can then be modeled as geometrically preserving or distorting the “shape” of the original experience. We applied our approach to data collected as participants watched and verbally recounted a television episode while undergoing functional neuroimaging. Participants’ recountings all preserved coarse spatial properties (essential narrative elements), but not fine spatial scale (low-level) details, of the episode’s trajectory. We also identified networks of brain structures sensitive to these trajectory shapes. Our work provides insights into how we preserve and distort our ongoing experiences when we encode them into episodic memories.

Introduction

What does it mean to remember something? In traditional episodic memory experiments (e.g., list-learning or trial-based experiments^{1,2}), remembering is often cast as a discrete, binary operation: each studied item may be separated from the rest of one’s experience and labeled as having been either recalled or forgotten. More nuanced studies might incorporate self-reported confidence measures as a proxy for memory strength, or ask participants to discriminate between recollecting the (contextual) details of an experience and having a general feeling of familiarity³. Using well-controlled, trial-based experimental designs, the field has amassed a wealth of information regarding human episodic memory⁴. However, there are fundamental properties of the external world and our memories that trial-based experiments are not well suited to capture^{5,6}. First, our experiences and memories are continuous, rather than discrete—isolating a naturalistic event from the context in which it occurs can substantially change its meaning. Second, whether or not the rememberer has precisely reproduced a specific set of words in describing a given experience is nearly orthogonal to how well they were actually able to remember it. In classic (e.g., list-learning)

memory studies, by contrast, the number or proportion of exact recalls is often considered to be a primary metric for assessing the quality of participants' memories. Third, one might remember the essence (or a general summary) of an experience but forget (or neglect to recount) particular low-level details. Capturing the essence of what happened is often a main goal of recounting an episodic memory to a listener, whereas the inclusion of specific low-level details is often less pertinent.

How might we formally characterize the "essence" of an experience, and whether it has been recovered by the rememberer? And how might we distinguish an experience's overarching essence from its low-level details? One approach is to start by considering some fundamental properties of the dynamics of our experiences. Each given moment of an experience tends to derive meaning from surrounding moments, as well as from longer-range temporal associations⁷⁻⁹. Therefore, the timecourse describing how an event unfolds is fundamental to its overall meaning. Further, this hierarchy formed by our subjective experiences at different timescales defines a context for each new moment^{10,11}, and plays an important role in how we interpret that moment and remember it later^{9,12}. Our memory systems can leverage these associations to form predictions that help guide our behaviors¹³. For example, as we navigate the world, the features of our subjective experiences tend to change gradually (e.g., the room or situation we find ourselves in at any given moment is strongly temporally autocorrelated), allowing us to form stable estimates of our current situation and behave accordingly^{14,15}.

Occasionally, this gradual drift of our ongoing experience is punctuated by sudden changes, or shifts (e.g., when we walk through a doorway¹⁶). Prior research suggests that these sharp transitions (termed "event boundaries") help to discretize our experiences (and their mental representations) into "events"¹⁶⁻²¹. The interplay between the stable (within-event) and transient (across-event) temporal dynamics of an experience also provides a potential framework for transforming experiences into memories that distills those experiences down to their essences. For example, prior work has shown that event boundaries can influence how we learn sequences of items^{18,21}, navigate¹⁷, and remember and understand narratives^{15,20}. This work also suggests a means of distinguishing the essence of an experience from its low-level details: The overall struc-

60 ture of events and event transitions reflects how the high-level experience unfolds (i.e., its essence),
61 while subtler event-level properties reflect its low-level details. Prior research has also implicated a
62 network of brain regions (including the hippocampus and the medial prefrontal cortex) in playing
63 a critical role in transforming experiences into structured and consolidated memories ²².

64 Here, we sought to examine how the temporal dynamics of a naturalistic experience were later
65 reflected in participants’ memories. We also sought to leverage the above conceptual insights
66 into the distinctions between an experience’s essence and its low-level details to build models
67 that explicitly quantified these distinctions. We analyzed an open dataset that comprised behav-
68 ioral and functional Magnetic Resonance Imaging (fMRI) data collected as participants viewed
69 and then verbally recounted an episode of the BBC television show *Sherlock*²³. We developed a
70 computational framework for characterizing the temporal dynamics of the moment-by-moment
71 content of the episode and of participants’ verbal recalls. Our framework uses topic modeling²⁴
72 to characterize the thematic conceptual (semantic) content present in each moment of the episode
73 and recalls by projecting each moment into a word embedding space. We then use hidden Markov
74 models^{25,26} to discretize this evolving semantic content into events. In this way, we cast both nat-
75 uralistic experiences and memories of those experiences as geometric “trajectories” through word
76 embedding space that describe how they evolve over time. Under this framework, successful
77 remembering entails verbally traversing the content trajectory of the episode, thereby reproducing
78 the shape (essence) of the original experience. Our framework captures the episode’s essence in
79 the sequence of geometric coordinates for its events, and its low-level details by examining its
80 within-event geometric properties.

81 Comparing the overall shapes of the topic trajectories for the episode and participants’ recalls
82 reveals which aspects of the episode’s essence were preserved (or lost) in the translation into mem-
83 ory. We also develop two metrics for assessing participants’ memories for low-level details: (1) the
84 “precision” with which a participant recounts details about each event, and (2) the “distinctive-
85 ness” of their recall for each event, relative to other events. We examine how these metrics relate to
86 overall memory performance as judged by third-party human annotators. We also compare and
87 contrast our general approach to studying memory for naturalistic experiences with standard met-

rics for assessing performance on more traditional memory tasks, such as list-learning. Last, we leverage our framework to identify networks of brain structures whose responses (as participants watched the episode) reflected the temporal dynamics of the episode and/or how participants would later recount it.

Results

To characterize the dynamic content of the *Sherlock* episode and participants’ subsequent recountings, we used a topic model²⁴ to discover the episode’s latent themes. Topic models take as inputs a vocabulary of words to consider and a collection of text documents, and return two output matrices. The first of these is a “topics matrix” whose rows are “topics” (or latent themes) and whose columns correspond to words in the vocabulary. The entries in the topics matrix reflect how each word in the vocabulary is weighted by each discovered topic. For example, a detective-themed topic might weight heavily on words like “crime,” and “search.” The second output is a “topic proportions matrix” with one row per document and one column per topic. The topic proportions matrix describes the mixture of discovered topics reflected in each document.

Chen et al. (2017) collected hand-annotated information about each of 1000 (manually delineated) time segments spanning the roughly 50 minute video used in their study²³. Each annotation included: a brief narrative description of what was happening, the location where the action took place, the names of any characters on the screen, and other similar details (for a full list of annotated features, see *Methods*). We took the union of all unique words (excluding stop words, such as “and,” “or,” “but,” etc.) across all features from all annotations as the vocabulary for the topic model. We then concatenated the sets of words across all features contained in overlapping sliding windows of (up to) 50 annotations, and treated each window as a single document for the purpose of fitting the topic model. Next, we fit a topic model with (up to) $K = 100$ topics to this collection of documents. We found that 32 unique topics (with non-zero weights) were sufficient to describe the time-varying content of the episode (see *Methods*; Fig. 1, Supp. Fig. 2). We note that our approach is similar in some respects to Dynamic Topic Models²⁷ in that we sought to characterize how the

114 thematic content of the episode evolved over time. However, whereas Dynamic Topic Models
115 are designed to characterize how the properties of collections of documents change over time,
116 our sliding window approach allows us to examine the topic dynamics within a single document
117 (or video). Specifically, our approach yielded (via the topic proportions matrix) a single “topic
118 vector” for each sliding window of annotations transformed by the topic model. We then stretched
119 (interpolated) the resulting windows-by-topics matrix to match the time series of the 1976 fMRI
120 volumes collected as participants viewed the episode.

121 The 32 topics we found were heavily character-focused (i.e., the top-weighted word in each
122 topic was nearly always a character) and could be roughly divided into themes centered around
123 Sherlock Holmes (the titular character), John Watson (Sherlock’s close confidant and assistant),
124 supporting characters (e.g., Inspector Lestrade, Sergeant Donovan, or Sherlock’s brother Mycroft),
125 or the interactions between various groupings of these characters (Supp. Fig. 2). This likely follows
126 from the frequency with which these terms appeared in the episode annotations. Several of
127 the identified topics were highly similar, which we hypothesized might allow us to distinguish
128 between subtle narrative differences if the distinctions between those overlapping topics were
129 meaningful. The topic vectors for each timepoint were also sparse, in that only a small number
130 of topics (typically one or two) tended to be “active” in any given timepoint (Fig. 2A). Further,
131 the dynamics of the topic activations appeared to exhibit persistence (i.e., given that a topic was
132 active in one timepoint, it was likely to be active in the following timepoint) along with occasional
133 rapid changes (i.e., occasionally topic weights would change abruptly from one timepoint to the
134 next). These two properties of the topic dynamics may be seen in the block diagonal structure of
135 the timepoint-by-timepoint correlation matrix (Fig. 2B) and reflect the gradual drift and sudden
136 shifts fundamental to the temporal dynamics of many real-world experiences, as well as television
137 episodes. Given this observation, we adapted an approach devised by Baldassano et al. (2017)²⁶,
138 and used a hidden Markov model (HMM) to identify the “event boundaries” where the topic
139 activations changed rapidly (i.e., the boundaries of the blocks in the temporal correlation matrix;
140 event boundaries identified by the HMM are outlined in yellow in Fig. 2B). Part of our model
141 fitting procedure required selecting an appropriate number of events into which the topic trajectory

142 should be segmented. To accomplish this, we used an optimization procedure that maximized
143 the difference between the topic weights for timepoints within an event versus timepoints across
144 multiple events (see *Methods*). We then created a stable summary of the content within each episode
145 event by averaging the topic vectors across the timepoints spanned by each event (Fig. 2C).

146 Given that the time-varying content of the episode could be segmented cleanly into discrete
147 events, we wondered whether participants' recalls of the episode also displayed a similar structure.
148 We applied the same topic model (already trained on the episode annotations) to each participant's
149 recalls. Analogously to how we parsed the time-varying content of the episode, to obtain similar
150 estimates for each participant's recall transcript, we treated each overlapping window of (up to)
151 10 sentences from their transcript as a document, and computed the most probable mix of topics
152 reflected in each timepoint's sentences. This yielded, for each participant, a number-of-windows
153 by number-of-topics topic proportions matrix that characterized how the topics identified in the
154 original episode were reflected in the participant's recalls. An important feature of our approach
155 is that it allows us to compare participants' recalls to events from the original episode, despite
156 that different participants used widely varying language to describe the events, and that those
157 descriptions often diverged in content, quality, and quantity from the episode annotations. This
158 ability to match up conceptually related text that differs in specific vocabulary, detail, and length
159 is an important benefit of projecting the episode and recalls into a shared topic space. An example
160 topic proportions matrix from one participant's recalls is shown in Figure 2D.

161 Although the example participant's recall topic proportions matrix has some visual similarity
162 to the episode topic proportions matrix, the time-varying topic proportions for the example par-
163 ticipant's recalls are not as sparse as those for the episode (compare Figs. 2A and D). Similarly,
164 although there do appear to be periods of stability in the recall topic dynamics (i.e., most topics are
165 active or inactive over contiguous blocks of time), the changes in topic activations that define event
166 boundaries appear less clearly delineated in participants' recalls than in the episode's annotations.
167 To examine these patterns in detail, we computed the timepoint-by-timepoint correlation matrix
168 for the example participant's recall topic proportions matrix (Fig. 2E). As in the episode correlation
169 matrix (Fig. 2B), the example participant's recall correlation matrix has a strong block diagonal

170 structure, indicating that their recalls are discretized into separated events. We used the same
171 HMM-based optimization procedure that we had applied to the episode’s topic proportions ma-
172 trix (see *Methods*) to estimate an analogous set of event boundaries in the participant’s recounting
173 of the episode (outlined in yellow). We carried out this analysis on all 17 participants’ recall topic
174 proportions matrices (Supp. Fig. 5).

175 Two clear patterns emerged from this set of analyses. First, although every individual partic-
176 ipant’s recalls could be segmented into discrete events (i.e., every individual participant’s recall
177 correlation matrix exhibited clear block diagonal structure; Supp. Fig. 5), each participant appeared
178 to have a unique “recall resolution,” reflected in the sizes of those blocks. While some participants’
179 recall topic proportions segmented into just a few events (e.g., Participants P4, P5, and P7), others’
180 segmented into many shorter-duration events (e.g., Participants P12, P13, and P17). This suggests
181 that different participants may be recalling the episode with different levels of detail—i.e., some
182 might recount only high-level essential plot details, whereas others might recount low-level details
183 instead (or in addition). The second clear pattern present in every individual participant’s recall
184 correlation matrix was that, unlike in the episode correlation matrix, there were substantial off-
185 diagonal correlations. One potential explanation for this finding is that the topic models, trained
186 only on episode annotations, do not capture topic proportions in participants’ “held-out” recalls
187 as accurately. A second possibility is that, whereas each event in the original episode was (largely)
188 separable from the others (Fig. 2B), in transforming those separable events into memory, partici-
189 pants appeared to be integrating across multiple events, blending elements of previously recalled
190 and not-yet-recalled content into each newly recalled event (Fig. 2E, Supp. Fig. 5)^{8,28,29}.

191 The above results demonstrate that topic models capture the dynamic conceptual content of
192 the episode and participants’ recalls of the episode. Further, the episode and recalls exhibit event
193 boundaries that can be identified automatically using HMMs to segment the dynamic content.
194 Next, we asked whether some correspondence might be made between the specific content of the
195 events the participants experienced while viewing the episode, and the events they later recalled.
196 We labeled each recall event as matching the episode event with the most similar (i.e., most highly
197 correlated) topic vector (Fig. 2G, Supp. Fig. 6). This yielded a sequence of “presented” events

198 from the original episode, and a (potentially differently ordered) sequence of “recalled” events
 199 for each participant. Analogous to classic list-learning studies, we can then examine participants’
 200 recall sequences by asking which events they tended to recall first (probability of first recall^{30–32};
 201 Fig. 3A); how participants most often transitioned between recalls of the events as a function of
 202 the temporal distance between them (lag-conditional response probability²; Fig. 3B); and which
 203 events they were likely to remember overall (serial position recall analyses¹; Fig. 3C). Some of
 204 the patterns we observed appeared to be similar to classic effects from the list-learning literature.
 205 For example, participants had a higher probability of initiating recall with early events (Fig. 3A)
 206 and a higher probability of transitioning to neighboring events with an asymmetric forward bias
 207 (Fig. 3B). However, unlike what is typically observed in list-learning studies, we did not observe
 208 patterns comparable to the primacy or recency serial position effects (Fig. 3C). We hypothesized
 209 that participants might be leveraging meaningful narrative associations and references over long
 210 timescales throughout the episode.

211 Clustering scores are often used by memory researchers to characterize how people organize
 212 their memories of words on a studied list³³. We defined analogous measures to characterize how
 213 participants organized their memories for episodic events (see *Methods* for details). Temporal
 214 clustering refers to the extent to which participants group their recall responses according to en-
 215 coding position. Overall, we found that sequentially viewed episode events tended to appear
 216 nearby in participants’ recall event sequences (mean clustering score: 0.732, SEM: 0.033). Par-
 217 ticipants with higher temporal clustering scores tended to exhibit better overall memory for the
 218 episode, according to both Chen et al. (2017)²³’s hand-counted numbers of recalled scenes from
 219 the episode (Pearson’s $r(15) = 0.49$, $p = 0.046$, 95% CI = [0.25, 0.76]) and the numbers of episode
 220 events that best-matched at least one recall event (i.e., model-estimated number of events recalled;
 221 Pearson’s $r(15) = 0.59$, $p = 0.013$, 95% CI = [0.31, 0.80]). Semantic clustering measures the extent
 222 to which participants cluster their recall responses according to semantic similarity³⁴. We found
 223 that participants tended to recall semantically similar episode events together (mean clustering
 224 score: 0.650, SEM: 0.032), and that semantic clustering scores were also related to both hand-
 225 counted (Pearson’s $r(15) = 0.65$, $p = 0.004$, 95% CI = [0.31, 0.85]) and model-estimated (Pearson’s

226 $r(15) = 0.58$, $p = 0.015$, 95% CI = [0.10, 0.83]) numbers of recalled events.

227 The above analyses illustrate how our framework for characterizing the dynamic conceptual
228 content of naturalistic episodes enables us to carry out analyses that have traditionally been
229 applied to much simpler list-learning paradigms. However, perhaps the most interesting aspects
230 of memory for naturalistic episodes are those that have no list-learning analogs. The nuances of
231 how one's memory for an event might capture some details, yet distort or neglect others, is central
232 to how we use our memory systems in daily life. Yet when researchers study memory in highly
233 simplified paradigms, those nuances are not typically observable. We next developed two novel,
234 continuous metrics, termed "precision" and "distinctiveness," aimed at characterizing distortions
235 in the conceptual content of individual recall events, and the conceptual overlap between how
236 people described different events.

237 Precision is intended to capture the "completeness" of recall, or how fully the presented content
238 was recapitulated in a participant's recounting. We define a recall event's precision as the maximum
239 correlation between the topic proportions of that recall event and any episode event (Fig. 4). In
240 other words, given that a recall event best matches a particular episode event, more precisely
241 recalled events overlap more strongly with the conceptual content of the original episode event.
242 When a given event is assigned a blend of several topics, as is often the case (Fig. 2), a high precision
243 score requires recapitulating the relative topic proportions during recall.

244 Distinctiveness is intended to capture the "specificity" of recall. In other words, distinctiveness
245 quantifies the extent to which a given recall event reflects the most similar episode event over
246 and above other episode events. Intuitively, distinctiveness is like a normalized variant of our
247 precision metric. Whereas precision solely measures how much detail about an event was captured
248 in someone's recall, distinctiveness penalizes details that also pertain to other episode events. We
249 define the distinctiveness of an event's recall as its precision expressed in standard deviation
250 units with respect to other episode events. Specifically, for a given recall event, we compute the
251 correlation between its topic vector and that of each episode event. This yields a distribution of
252 correlation coefficients (one per episode event). We subtract the mean and divide by the standard
253 deviation of this distribution to z-score the coefficients. The maximum value in this distribution

(which, by definition, belongs to the episode event that best matches the given recall event) is that recall event’s distinctiveness score. In this way, recall events that match one episode event far better than all other episode events will receive a high distinctiveness score. By contrast, a recall event that matches all episode events roughly equally will receive a comparatively low distinctiveness score.

In addition to examining how precisely and distinctively participants recalled individual events, one may also use these metrics to summarize each participant’s performance by averaging across a participant’s event-wise precision or distinctiveness scores. This enables us to quantify how precisely a participant tended to recall subtle within-event details, as well as how specific (distinctive) those details were to individual events from the episode. Participants’ average precision and distinctiveness scores were strongly correlated ($r(15) = 0.90$, $p < 0.001$, 95% CI = [0.66,0.96]). This indicates that participants who tended to precisely recount low-level details of episode events also tended to do so in an event-specific way (e.g., as opposed to detailing recurring themes that were present in most or all episode events; this behavior would have resulted in high precision but low distinctiveness). We found that, across participants, higher precision scores were positively correlated with the numbers of both model-estimated events ($r(15) = 0.90$, $p < 0.001$, 95% CI = [0.54,0.96]) and hand-annotated scenes ($r(15) = 0.60$, $p = 0.010$, 95% CI = [0.02,0.83]) that participants recalled. Participants’ average distinctiveness scores were also correlated with their numbers of model-estimated recalled events ($r(15) = 0.71$, $p = 0.001$, 95% CI = [-0.07,0.90]) and marginally significantly correlated with their numbers of hand-annotated ($r(15) = 0.45$, $p = 0.068$, 95% CI = [-0.21,0.79]).

Examining individual recalls of the same episode event can provide insights into how the above precision and distinctiveness scores may be used to characterize similarities and differences in how different people describe the same shared experience. In Figure 5, we compare recalls for the same episode event from the participants with the highest (P17) and lowest (P6) precision scores. From the HMM-identified episode event boundaries, we recovered the set of annotations describing the content of a single episode event (event 21; Fig. 5C), and divided them into different color-coded sections for each action or feature described. Next, we used an analogous approach to identify

the set of sentences comprising the corresponding recall event from each of the two example participants (Fig. 5D). We then colored all words describing actions and features in the transcripts shown in Panel D according to the color-coded annotations in Panel C. Visual comparison of these example recalls reveals that the more precise recall captures more of the episode event’s content, and in greater detail.

Figure 5 also illustrates the differences between high and low distinctiveness scores. We extracted the set of sentences comprising the most distinctive recall event (P9) and least distinctive recall event (P6) corresponding to the example episode event shown in Panel C (event 21). We also extracted the annotations for all episode events whose content these participants’ single recall events touched on. We assigned each episode event a unique color (Fig. 5E), and colored each recalled sentence (Panel F) according to the episode events they best matched. Visual inspection of Panel F reveals that the most distinctive recall’s content is tightly concentrated around event 21, whereas the least distinctive recall incorporates content from a much wider range of episode events.

The preceding analyses sought to characterize how participants’ recountings of individual episode events captured the low-level details of each event. Next, we sought to characterize how participants’ recountings of the full episode captured its high-level essence—i.e., the shape of the episode’s trajectory through word embedding (topic) space. To visualize the essence of the episode and each participant’s recall trajectory³⁵, we projected the topic proportions matrices for the episode and recalls onto a shared two-dimensional space using Uniform Manifold Approximation and Projection (UMAP)³⁶. In this lower-dimensional space, each point represents a single episode or recall event, and the distances between the points reflect the distances between the events’ associated topic vectors (Fig. 6). In other words, events that are nearer to each other in this space are more semantically similar, and those that are farther apart are less so.

Visual inspection of the episode and recall topic trajectories reveals a striking pattern. First, the topic trajectory of the episode (which reflects its dynamic content; Fig. 6A) is captured nearly perfectly by the averaged topic trajectories of participants’ recalls (Fig. 6B). To assess the consistency of these recall trajectories across participants, we asked: given that a participant’s recall trajectory

310 had entered a particular location in the reduced topic space, could the position of their next recalled
 311 event be predicted reliably? For each location in the reduced topic space, we computed the set of
 312 line segments connecting successively recalled events (across all participants) that intersected that
 313 location (see *Methods*, Supp. Fig. 4). We then computed (for each location) the distribution of angles
 314 formed by the lines defined by those line segments and a fixed reference line (the x -axis). Rayleigh
 315 tests revealed the set of locations in topic space at which these across-participant distributions
 316 exhibited reliable peaks (blue arrows in Fig. 6B reflect significant peaks at $p < 0.05$, corrected). We
 317 observed that the locations traversed by nearly the entire episode trajectory exhibited such peaks.
 318 In other words, participants' recalls exhibited similar trajectories to each other that also matched the
 319 trajectory of the original episode (Fig. 6C). This is especially notable when considering the fact that
 320 the number of HMM-identified recall events (dots in Fig. 6C) varied considerably across people,
 321 and that every participant used different words to describe what they had remembered happening
 322 in the episode. Differences in the numbers of recall events appear in participants' trajectories
 323 as differences in the sampling resolution along the trajectory. We note that this framework also
 324 provides a means of disentangling classic "proportion recalled" measures (i.e., the proportion of
 325 episode events described in participants' recalls) from participants' abilities to recapitulate the
 326 episode's essence (i.e., the similarity between the shapes of the original episode trajectory and that
 327 defined by each participant's recounting of the episode).

328 In addition to enabling us to visualize the episode's high-level essence, describing the episode
 329 as a geometric trajectory also enables us to drill down to individual words and quantify how each
 330 word relates to the memorability of each event. This provides another approach to examining
 331 participants' recall for low-level details beyond the precision and distinctiveness measures we
 332 defined above. The results displayed in Figures 3C and 5A suggest that certain events were
 333 remembered better than others. Given this, we next asked whether the events that were generally
 334 remembered precisely or imprecisely tended to reflect particular content. Because our analysis
 335 framework projects the dynamic episode content and participants' recalls into a shared space, and
 336 because the dimensions of that space represent topics (which are, in turn, sets of weights over
 337 known words in the vocabulary), we are able to recover the weighted combination of words that

338 make up any point (i.e., topic vector) in this space. We first computed the average precision with
339 which participants recalled each of the 30 episode events (Fig. 7A; note that this result is analogous
340 to a serial position curve created from our precision metric). We then computed a weighted average
341 of the topic vectors for each episode event, where the weights reflected how precisely each event
342 was recalled. To visualize the result, we created a “wordle” image³⁷ where words weighted more
343 heavily by more precisely remembered topics appear in a larger font (Fig. 7B, green box). Across
344 the full episode, content that weighted heavily on topics and words central to the major foci of
345 the episode (e.g., the names of the two main characters, “Sherlock” and “John,” and the address
346 of a major recurring location, “221B Baker Street”) was best remembered. An analogous analysis
347 revealed which themes were less-precisely remembered. Here, in computing the weighted average
348 over events’ topic vectors, we weighted each event in inverse proportion to its average precision
349 (Fig. 7B, red box). The least precisely remembered episode content reflected information that was
350 extraneous to the episode’s essence, such as the proper names of relatively minor characters (e.g.,
351 “Mike,” “Molly,” and “Lestrade”) and locations (e.g., “St. Bartholomew’s Hospital”).

352 A similar result emerged from assessing the topic vectors for individual episode and recall
353 events (Fig. 7C). Here, for each of the three most and least precisely remembered episode events, we
354 have constructed two wordles: one from the original episode event’s topic vector (left) and a second
355 from the average recall topic vector for that event (right). The three most precisely remembered
356 events (circled in green) correspond to scenes integral to the central plot-line: a mysterious figure
357 spying on John in a phone booth; John meeting Sherlock at Baker St. to discuss the murders;
358 and Sherlock laying a trap to catch the killer. Meanwhile, the least precisely remembered events
359 (circled in red) reflect scenes that comprise minor plot points: a video of singing cartoon characters
360 that participants viewed in an introductory clip prior to the main episode; John asking Molly
361 about Sherlock’s habit of over-analyzing people; and Sherlock noticing evidence of Anderson’s
362 and Donovan’s affair.

363 The results this far inform us about which aspects of the dynamic content in the episode partici-
364 pants watched were preserved or altered in participants’ memories. We next carried out a series of
365 analyses aimed at understanding which brain structures might facilitate these preservations and

transformations between the participants' shared experience of watching the episode and their subsequent memories of the episode. In the first analysis, we sought to identify brain structures that were sensitive to the dynamic unfolding of the episode's content, as characterized by its topic trajectory. We used a searchlight procedure to identify clusters of voxels whose activity patterns displayed a proximal temporal correlation structure (as participants watched the episode) matching that of the original episode's topic proportions (Fig. 8A; see *Methods* for additional details). In a second analysis, we sought to identify brain structures whose responses (during episode viewing) reflected how each participant would later structure their recounting of the episode. We used a searchlight procedure to identify clusters of voxels whose proximal temporal correlation matrices matched that of the topic proportions matrix for each participant's recall transcript (Figs. 8B; see *Methods* for additional details). To ensure our searchlight procedure identified regions specifically sensitive to the temporal structure of the episode or recalls (i.e., rather than those with a temporal autocorrelation length similar to that of the episode and recalls), we performed a phase shift-based permutation correction (see *Methods*). As shown in Figure 8C, the episode-driven searchlight analysis revealed a distributed network of regions that may play a role in processing information relevant to the narrative structure of the episode. The recall-driven searchlight analysis revealed a second network of regions (Fig. 8D) that may facilitate a person-specific transformation of one's experience into memory. In identifying regions whose responses to ongoing experiences reflect how those experiences will be remembered later, this latter analysis extends classic "subsequent memory effect analyses"³⁸ to the domain of naturalistic experiences.

The searchlight analyses described above yielded two distributed networks of brain regions whose activity timecourses tracked with the temporal structure of the episode (Fig. 8C) or participants' subsequent recalls (Fig. 8D). We next sought to gain greater insight into the structures and functional networks our results reflected. To accomplish this, we performed an additional, exploratory analysis using Neurosynth³⁹. Given an arbitrary statistical map as input, Neurosynth performs a massive automated meta-analysis, returning a frequency-ranked list of terms used in neuroimaging papers that report similar statistical maps. We ran Neurosynth on the (unthresholded) permutation-corrected maps for the episode- and recall-driven searchlight analyses. The

394 top ten terms with maximally similar meta-analysis images identified by Neurosynth are shown
395 in Figure 8.

396 Discussion

397 Explicitly modeling the dynamic content of a naturalistic stimulus and participants' memories
398 enabled us to connect the present study of naturalistic recall with an extensive prior literature that
399 has used list-learning paradigms to study memory⁴, as in Figure 3. We found some similarities
400 between how participants in the present study recounted a television episode and how participants
401 typically recall memorized random word lists. However, our broader claim is that word lists miss
402 out on fundamental aspects of naturalistic memory more like the sort of memory we rely on
403 in everyday life. For example, there are no random word list analogs of character interactions,
404 conceptual dependencies between temporally distant episode events, the sense of solving a mystery
405 that pervades the *Sherlock* episode, or the myriad other features of the episode that convey deep
406 meaning and capture interest. Nevertheless, each of these properties affects how people process
407 and engage with the episode as they are watching it, and how they remember it later. The
408 overarching goal of the present study is to characterize how the rich dynamics of the episode affect
409 the rich behavioral and neural dynamics of how people remember it.

410 Our work casts remembering as reproducing (behaviorally and neurally) the topic trajectory,
411 or "shape," of an experience, thereby drawing implicit analogies between mentally navigating
412 through word embedding spaces and physically navigating through spatial environments^{40–42}.
413 When we characterized memory for a television episode using this framework, we found that
414 every participant's recounting of the episode recapitulated the low spatial frequency details of
415 the shape of its trajectory through topic space (Fig. 6). We termed this narrative scaffolding the
416 episode's essence. Where participants' behaviors varied most was in their tendencies to recount
417 specific low-level details from each episode event. Geometrically, this appears as high spatial
418 frequency distortions in participants' recall trajectories relative to the trajectory of the original
419 episode (Fig. 7). We developed metrics to characterize the precision (recovery of any and all event-

level information) and distinctiveness (recovery of event-specific information). We also used word cloud visualizations to interpret the details of these event-level distortions.

The neural analyses we carried out (Fig. 8) also leveraged our geometric framework for characterizing the shapes of the episode and participants' recountings. We identified one network of regions whose responses tracked with temporal correlations in the conceptual content of the episode (as quantified by topic models applied to a set of annotations about the episode). This network included orbitofrontal cortex, ventromedial prefrontal cortex, and striatum, among others. As reviewed by Ranganath and Ritchey (2012)¹³, several of these regions are members of the "anterior temporal system," which has been implicated in assessing and processing the familiarity of ongoing experiences, emotions, social cognition, and reward. A second network we identified tracked with temporal correlations in the idiosyncratic conceptual content of participants' subsequent recountings of the episode. This network included occipital cortex, extrastriate cortex, fusiform gyrus, and the precuneus. Several of these regions are members of the "posterior medial system"¹³, which has been implicated in matching incoming cues about the current situation to internally maintained "situation models" that specify the parameters and expectations inherent to the current situation^{14,15}. Taken together, our results support the notion that these two (partially overlapping) networks work in coordination to make sense of our ongoing experiences, distort them in a way that links them with our prior knowledge and experiences, and encodes those distorted representations into memory for our later use. Our work also provides a potential framework for modeling and elucidating "memory schemas"—i.e., cognitive abstractions that may be applied to multiple related experiences^{43,44}. For example, the event-level geometric scaffolding of an experience (e.g., Fig. 6A) might reflect its underlying schema, and experiences that share similar schemas might have similar shapes. This could also help explain how brain structures including the ventromedial prefrontal cortex⁴³ (Fig. 8) might acquire or apply schema knowledge across different experiences (i.e., by learning patterns in the schema's shape).

Our general approach draws inspiration from prior work aimed at elucidating the neural and behavioral underpinnings of how we process dynamic naturalistic experiences and remember them later. Our approach to identifying neural responses to naturalistic stimuli (including experiences)

entails building an explicit model of the stimulus dynamics and searching for brain regions whose responses are consistent with the model^{45,46}. Building an explicit model of these dynamics also enables us to match up different people’s recountings of a common shared experience, despite individual differences⁴⁷. In prior work, a series of studies from Uri Hasson’s group^{7,23,26,48,49} have presented a clever alternative approach: rather than building an explicit stimulus model, these studies instead search for brain responses to the stimulus that are reliably similar across individuals. So called “inter-subject correlation” (ISC) and “inter-subject functional connectivity” (ISFC) analyses effectively treat other people’s brain responses to the stimulus as a “model” of how its features change over time⁵⁰. These purely brain-driven approaches are well suited to identifying which brain structures exhibit similar stimulus-driven responses across individuals. Further, because neural response dynamics are observed data (rather than model approximations), such approaches do not require a detailed understanding of which stimulus properties or features might be driving the observed responses. However, this also means that the specific stimulus features driving those responses are typically opaque to the researcher. Our approach is complementary. By explicitly modeling the stimulus dynamics, we are able to relate specific stimulus features to behavioral and neural dynamics. However, when our model fails to accurately capture the stimulus dynamics that are truly driving behavioral and neural responses, our approach necessarily yields an incomplete characterization of the neural basis of the processes we are studying.

Other recent work has used HMMs to discover latent event structure in neural responses to naturalistic stimuli²⁶. By applying HMMs to our explicit models of stimulus and memory dynamics, we gain a more direct understanding of those state dynamics. For example, we found that although the events comprising each participant’s recalls recapitulated the episode’s essence, participants differed in the resolution of their recounting of low-level details. In turn, these individual behavioral differences were reflected in differences in neural activity dynamics as participants watched the television episode.

Our approach also draws inspiration from the growing field of word embedding models. The topic models²⁴ we used to embed text from the episode annotations and participants’ recall transcripts are just one of many models that have been studied in an extensive literature. The earliest

approaches to word embedding, including latent semantic analysis⁵¹, used word co-occurrence statistics (i.e., how often pairs of words occur in the same documents contained in the corpus) to derive a unique feature vector for each word. The feature vectors are constructed so that words that co-occur more frequently have feature vectors that are closer (in Euclidean distance). Topic models are essentially an extension of those early models, in that they attempt to explicitly model the underlying causes of word co-occurrences by automatically identifying the set of themes or topics reflected across the documents in the corpus. More recent work on these types of semantic models, including word2vec⁵², the Universal Sentence Encoder⁵³, and Generative Pre-trained Transformers (e.g., GPT-2⁵⁴ and GTP-3⁵⁵) use deep neural networks to attempt to identify the deeper conceptual representations underlying each word. Despite the growing popularity of these sophisticated deep learning-based embedding models, we chose to prioritize interpretability of the embedding dimensions (e.g., Fig. 7) over raw performance (e.g., with respect to some pre-defined benchmark). Nevertheless, we note that our general framework is, in principle, robust to the specific choice of language model as well as other aspects of our computational pipeline. For example, the word embedding model, timeseries segmentation model, and the episode-recall matching function could each be customized to suit a particular question space or application. Indeed, for some questions, interpretability of the embeddings may not be a priority, and thus other text embedding approaches (including the deep learning-based models described above) may be preferable. Further work will be needed to explore the influence of particular models on our framework’s predictions and performance.

Speculatively, our work may have broad implications for how we characterize and assess memory in real-world settings, such as the classroom or physician’s office. For example, the most commonly used classroom evaluation tools involve simply computing the proportion of correctly answered exam questions. Our work suggests that this approach is only loosely related to what educators might really want to measure: how well did the students understand the key ideas presented in the course? Under this typical framework of assessment, the same exam score of 50% could be ascribed to two very different students: one who attended to the full course but struggled to learn more than a broad overview of the material, and one who attended to only half of the

504 course but understood the attended material perfectly. Instead, one could apply our computational
505 framework to build explicit dynamic content models of the course material and exam questions.
506 This approach might provide a more nuanced and specific view into which aspects of the material
507 students had learned well (or poorly). In clinical settings, memory measures that incorporate such
508 explicit content models might also provide more direct evaluations of patients' memories, and of
509 doctor-patient interactions.

510 **Methods**

511 **Paradigm and data collection**

512 Data were collected by Chen et al. (2017)²³. In brief, participants ($n = 22$) viewed the first 48
513 minutes of "A Study in Pink," the first episode of the BBC television show *Sherlock*, while fMRI
514 volumes were collected (TR = 1500 ms). Participants were pre-screened to ensure they had never
515 seen any episode of the show before. The stimulus was divided into a 23 min (946 TR) and a
516 25 min (1030 TR) segment to mitigate technical issues related to the scanner. After finishing the
517 clip, participants were instructed to "describe what they recalled of the [episode] in as much detail
518 as they could, to try to recount events in the original order they were viewed in, and to speak for at
519 least 10 minutes if possible but that longer was better. They were told that completeness and detail
520 were more important than temporal order, and that if at any point they realized they had missed
521 something, to return to it. Participants were then allowed to speak for as long as they wished, and
522 verbally indicated when they were finished (e.g., 'I'm done')." ²³ Five participants were dropped
523 from the original dataset due to excessive head motion (2 participants), insufficient recall length (2
524 participants), or falling asleep during stimulus viewing (1 participant), resulting in a final sample
525 size of $n = 17$. For additional details about the testing procedures and scanning parameters, see
526 Chen et al. (2017)²³. The testing protocol was approved by Princeton University's Institutional
527 Review Board.

528 After preprocessing the fMRI data and warping the images into a standard (3 mm³ MNI) space,

the voxel activations were z-scored (within voxel) and spatially smoothed using a 6 mm (full width at half maximum) Gaussian kernel. The fMRI data were also cropped so that all episode-viewing data were aligned across participants. This included a constant 3 TR (4.5 s) shift to account for the lag in the hemodynamic response. All of these preprocessing steps followed Chen et al. (2017)²³, where additional details may be found.

The video stimulus was divided into 1000 fine-grained “time segments” and annotated by an independent coder. For each of these 1000 annotations, the following information was recorded: a brief narrative description of what was happening, the location where the time segment took place, whether that location was indoors or outdoors, the names of all characters on-screen, the name(s) of the character(s) in focus in the shot, the name(s) of the character(s) currently speaking, the camera angle of the shot, a transcription of any text appearing on-screen, and whether or not there was music present in the background. Each time segment was also tagged with its onset and offset time, in both seconds and TRs.

Statistics

All statistical tests performed in the behavioral analyses were two-sided. All statistical tests performed in the neural data analyses were two-sided, except for the permutation-based thresholding, which was one-sided. In this case, we were specifically interested in identifying voxels whose activation time series reflected the temporal structure of the episode and recall topic proportions matrices to a greater extent than that of the phase-shifted matrices. The 95% confidence intervals we reported for each correlation were estimated by generating 10000 “bootstrap” distributions of correlation coefficients by sampling (with replacement) from the observed data.

Modeling the dynamic content of the episode and recall transcripts

Topic modeling

The input to the topic model we trained to characterize the dynamic content of the episode comprised 998 hand-generated annotations of short (mean: 2.96s) time segments spanning the video

clip (Chen et al., 2017²³ generated 1000 annotations total; we removed two annotations referring to a break between the first and second scan sessions, during which no fMRI data were collected). We concatenated the text for all of the annotated features within each segment, creating a “bag of words” describing its content, and performed some minor preprocessing (e.g., stemming possessive nouns and removing punctuation). We then re-organized the text descriptions into overlapping sliding windows spanning (up to) 50 annotations each. In other words, we estimated the “context” for each annotated segment using the text descriptions of the preceding 25 annotations, the present annotations, and the following 24 annotations. To model the context for annotations near the beginning of the episode (i.e., within 25 of the beginning or end), we created overlapping sliding windows that grew in size from one annotation to the full length. We also tapered the sliding window lengths at the end of the episode, whereby time segments within fewer than 24 annotations of the end of the episode were assigned sliding windows that extended to the end of the episode. This procedure ensured that each annotation’s content was represented in the text corpus an equal number of times.

We trained our model using these overlapping text samples with `scikit-learn` version 0.19.1⁵⁶, called from our high-dimensional visualization and text analysis software, `HyperTools`³⁵. Specifically, we used the `CountVectorizer` class to transform the text from each window into a vector of word counts (using the union of all words across all annotations as the “vocabulary,” excluding English stop words); this yielded a number-of-windows by number-of-words “word count” matrix. We then used the `LatentDirichletAllocation` class (topics=100, method=‘batch’) to fit a topic model²⁴ to the word count matrix, yielding a number-of-windows (1047) by number-of-topics (100) “topic proportions” matrix. The topic proportions matrix describes the gradually evolving mix of topics (latent themes) present in each annotated time segment of the episode. Next, we transformed the topic proportions matrix to match the 1976 fMRI volume acquisition times. We assigned each topic vector to the timepoint (in seconds) midway between the beginning of the first annotation and the end of the last annotation in its corresponding sliding text window. By doing so, we warped the linear temporal distance between consecutive topic vectors to align with the inconsistent temporal distance between consecutive annotations (whose durations varied greatly).

582 We then rescaled these timepoints to 1.5s TR units, and used linear interpolation to estimate a topic
583 vector for each TR. This resulted in a number-of-TRs (1976) by number-of-topics (100) matrix.

584 We created similar topic proportions matrices using hand-annotated transcripts of each par-
585 ticipant’s verbal recall of the episode²³. We tokenized the transcript into a list of sentences, and
586 then re-organized the list into overlapping sliding windows spanning (up to) 10 sentences each,
587 analogously to how we parsed the episode annotations. In turn, we transformed each window’s
588 sentences into a word count vector (using the same vocabulary as for the episode model), then
589 used the topic model already trained on the episode scenes to compute the most probable topic
590 proportions for each sliding window. This yielded a number-of-windows (range: 83–312) by
591 number-of-topics (100) topic proportions matrix for each participant. These reflected the dynamic
592 content of each participant’s recalls. For details on how we selected the episode and recall window
593 lengths and number of topics, see *Supplementary Information* and Supplementary Figure 1.

594 **Segmenting topic proportions matrices into discrete events using hidden Markov Models**

595 We parsed the topic proportions matrices of the episode and participants’ recalls into discrete
596 events using hidden Markov Models (HMMs)²⁵. Given the topic proportions matrix (describing
597 the mix of topics at each timepoint) and a number of states, K , an HMM recovers the set of state
598 transitions that segments the timeseries into K discrete states. Following Baldassano et al. (2017)²⁶,
599 we imposed an additional set of constraints on the discovered state transitions that ensured that
600 each state was encountered exactly once (i.e., never repeated). We used the BrainIAK toolbox⁵⁷ to
601 implement this segmentation.

602 We used an optimization procedure to select the appropriate K for each topic proportions matrix.
603 Prior studies on narrative structure and processing have shown that we both perceive and internally
604 represent the world around us at multiple, hierarchical timescales^{7,23,26,44,58,59}. However, for the
605 purposes of our framework, we sought to identify the single timeseries of event representations
606 that was emphasized most heavily in the temporal structure of the episode and of each participant’s
607 recall. We quantified this as the set of K states that maximized the similarity between topic vectors
608 for timepoints comprising each state, while minimizing the similarity between topic vectors for

timepoints across different states. Specifically, we computed (for each matrix)

$$\underset{K}{\operatorname{argmax}} [W_1(a, b)],$$

where a was the distribution of within-state topic vector correlations, and b was the distribution of across-state topic vector correlations. We computed the first Wasserstein distance (W_1 , also known as “Earth mover’s distance”^{60,61}) between these distributions for a large range of possible K -values (range [2, 50]), and selected the K that yielded the maximum value. Figure 2B displays the event boundaries returned for the episode, and Supplementary Figure 5 displays the event boundaries returned for each participant’s recalls. See Supplementary Figure 7 for the optimization functions for the episode and recalls. After obtaining these event boundaries, we created stable estimates of the content represented in each event by averaging the topic vectors across timepoints between each pair of event boundaries. This yielded a number-of-events by number-of-topics matrix for the episode and recalls from each participant.

Naturalistic extensions of classic list-learning analyses

In traditional list-learning experiments, participants view a list of items (e.g., words) and then recall the items later. Our episode-recall event matching approach affords us the ability to analyze memory in a similar way. The episode and recall events can be treated analogously to studied and recalled “items” in a list-learning study. We can then extend classic analyses of memory performance and dynamics (originally designed for list-learning experiments) to the more naturalistic episode recall task used in this study.

Perhaps the simplest and most widely used measure of memory performance is “accuracy”—i.e., the proportion of studied (experienced) items (in this case, episode events) that the participant later remembered. Chen et al. (2017)²³ used this method to rate each participant’s memory quality by computing the proportion of (50 manually identified) scenes mentioned in their recall. We found a strong across-participants correlation between these independent ratings and the proportion of 30 HMM-identified episode events matched to participants’ recalls (Pearson’s

633 $r(15) = 0.71, p = 0.002, 95\% \text{ CI} = [0.39, 0.88]$). We further considered a number of more nuanced
634 memory performance measures that are typically associated with list-learning studies. We also
635 provide a software package, *Quail*, for carrying out these analyses⁶².

636 **Probability of first recall (PFR).** PFR curves^{30–32} reflect the probability that an item will be
637 recalled first, as a function of its serial position during encoding. To carry out this analysis, we
638 initialized a number-of-participants (17) by number-of-episode-events (30) matrix of zeros. Then,
639 for each participant, we found the index of the episode event that was recalled first (i.e., the episode
640 event whose topic vector was most strongly correlated with that of the first recall event) and filled
641 in that index in the matrix with a 1. Finally, we averaged over the rows of the matrix, resulting in a
642 1 by 30 array representing the proportion of participants that recalled an event first, as a function
643 of the order of the event’s appearance in the episode (Fig. 3A).

644 **Lag conditional probability curve (lag-CRP).** The lag-CRP curve² reflects the probability of
645 recalling a given item after the just-recalled item, as a function of their relative encoding positions
646 (lag). In other words, a lag of 1 indicates that a recalled item was presented immediately after
647 the previously recalled item, and a lag of -3 indicates that a recalled item came 3 items before the
648 previously recalled item. For each recall transition (following the first recall), we computed the
649 lag between the current recall event and the next recall event, normalizing by the total number
650 of possible transitions. This yielded a number-of-participants (17) by number-of-lags (-29 to +29;
651 58 lags total excluding lags of 0) matrix. We averaged over the rows of this matrix to obtain a
652 group-averaged lag-CRP curve (Fig. 3B).

653 **Serial position curve (SPC).** SPCs¹ reflect the proportion of participants that remember each
654 item as a function of the item’s serial position during encoding. We initialized a number-of-
655 participants (17) by number-of-episode-events (30) matrix of zeros. Then, for each recalled event,
656 for each participant, we found the index of the episode event that the recalled event most closely
657 matched (via the correlation between the events’ topic vectors) and entered a 1 into that position
658 in the matrix. This resulted in a matrix whose entries indicated whether or not each event was

659 recalled by each participant (depending on whether the corresponding entires were set to one or
660 zero). Finally, we averaged over the rows of the matrix to yield a 1 by 30 array representing the
661 proportion of participants that recalled each event as a function of the events' order appearance in
662 the episode (Fig. 3C).

663 **Temporal clustering scores.** Temporal clustering describes a participant's tendency to organize
664 their recall sequences by the learned items' encoding positions. For instance, if a participant
665 recalled the episode events in the exact order they occurred (or in exact reverse order), this would
666 yield a score of 1. If a participant recalled the events in random order, this would yield an expected
667 score of 0.5. For each recall event transition (and separately for each participant), we sorted all
668 not-yet-recalled events according to their absolute lag (i.e., distance away in the episode). We
669 then computed the percentile rank of the next event the participant recalled. We averaged these
670 percentile ranks across all of the participant's recalls to obtain a single temporal clustering score
671 for the participant.

672 **Semantic clustering scores.** Semantic clustering describes a participant's tendency to recall se-
673 mantically similar presented items together in their recall sequences. Here, we used the topic
674 vectors for each event as a proxy for its semantic content. Thus, the similarity between the se-
675 mantic content for two events can be computed by correlating their respective topic vectors. For
676 each recall event transition, we sorted all not-yet-recalled events according to how correlated the
677 topic vector of the closest-matching episode event was to the topic vector of the closest-matching
678 episode event to the just-recalled event. We then computed the percentile rank of the observed
679 next recall. We averaged these percentile ranks across all of the participant's recalls to obtain a
680 single semantic clustering score for the participant.

681 **Averaging correlations**

682 In all instances where we performed statistical tests involving precision or distinctiveness scores
683 (Fig. 5), we used the Fisher z-transformation⁶³ to stabilize the variance across the distribution of

684 correlation values prior to performing the test. Similarly, when averaging precision or distinctive-
685 ness scores, we z-transformed the scores prior to computing the mean, and inverse z-transformed
686 the result.

687 **Visualizing the episode and recall topic trajectories**

688 We used the UMAP algorithm³⁶ to project the 100-dimensional topic space onto a two-dimensional
689 space for visualization (Figs. 6, 7). To ensure that all of the trajectories were projected onto the
690 same lower dimensional space, we computed the low-dimensional embedding on a “stacked”
691 matrix created by vertically concatenating the events-by-topics topic proportions matrices for the
692 episode, the across-participants average recalls and all 17 individual participants’ recalls. We then
693 separated the rows of the result (a total-number-of-events by two matrix) back into individual
694 matrices for the episode topic trajectory, the across-participant average recall trajectory, and the
695 trajectories for each individual participant’s recalls (Fig. 6). This general approach for discovering
696 a shared low-dimensional embedding for a collections of high-dimensional observations follows
697 our prior work on manifold learning³⁵.

698 We optimized the manifold space for visualization based on two criteria: First, that the 2D em-
699 bedding of the episode trajectory should reflect its original 100-dimensional structure as faithfully
700 as possible. Second, that the path traversed by the embedded episode trajectory should intersect
701 itself a minimal number of times. The first criteria helps bolster the validity of visual intuitions
702 about relationships between sections of episode content, based on their locations in the embed-
703 ding space. The second criteria was motivated by the observed low off-diagonal values in the
704 episode trajectory’s temporal correlation matrix (suggesting that the same topic-space coordinates
705 should not be revisited; see Fig. 2A). For further details on how we created this low-dimensional
706 embedding space, see *Supplementary Information*.

707 **Estimating the consistency of flow through topic space across participants**

708 In Figure 6B, we present an analysis aimed at characterizing locations in topic space that different
709 participants move through in a consistent way (via their recall topic trajectories; also see Supp.

Fig. 4). The two-dimensional topic space used in our visualizations (Fig. 6) comprised a 60×60 (arbitrary units) square. We tiled this space with a 50×50 grid of evenly spaced vertices, and defined a circular area centered on each vertex whose radius was two times the distance between adjacent vertices (i.e., 2.4 units). For each vertex, we examined the set of line segments formed by connecting each pair successively recalled events, across all participants, that passed through this circle. We computed the distribution of angles formed by those segments and the x -axis, and used a Rayleigh test to determine whether the distribution of angles was reliably “peaked” (i.e., consistent across all transitions that passed through that local portion of topic space). To create Figure 6B, we drew an arrow originating from each grid vertex, pointing in the direction of the average angle formed by the line segments that passed within 2.4 units. We set the arrow lengths to be inversely proportional to the p -values of the Rayleigh tests at each vertex. Specifically, for each vertex we converted all of the angles of segments that passed within 2.4 units to unit vectors, and we set the arrow lengths at each vertex proportional to the length of the (circular) mean vector. We also indicated any significant results ($p < 0.05$, corrected using the Benjamini-Hochberg procedure) by coloring the arrows in blue (darker blue denotes a lower p -value, i.e., a longer mean vector); all tests with $p \geq 0.05$ are displayed in gray and given a lower opacity value.

Searchlight fMRI analyses

In Figure 8, we present two analyses aimed at identifying brain regions whose responses (as participants viewed the episode) exhibited a particular temporal structure. We developed a searchlight analysis wherein we constructed a $5 \times 5 \times 5$ cube of voxels centered on each voxel in the brain²³, and for each of these cubes, computed the temporal correlation matrix of the voxel responses during episode viewing. Specifically, for each of the 1976 volumes collected during episode viewing, we correlated the activity patterns in the given cube with the activity patterns (in the same cube) collected during every other timepoint. This yielded a 1976×1976 correlation matrix for each cube. Note: participant 5’s scan ended 75s early, and in Chen et al. (2017)²³’s publicly released dataset, their scan data was zero-padded to match the length of the other participants’. For our searchlight analyses, we removed this padded data (i.e., the last 50 TRs), resulting in a 1925×1925 correlation

737 matrix for each cube in participant 5's brain.

738 Next, we constructed a series of "template" matrices. The first template reflected the time-
739 course of the episode's topic proportions matrix, and the others reflected the timecourse of each
740 participant's recall topic proportions matrix. To construct the episode template, we computed the
741 correlations between the topic proportions estimated for every pair of TRs (prior to segmenting
742 the topic proportions matrices into discrete events; i.e., the correlation matrix shown in Figs. 2B
743 and 8A). We constructed similar temporal correlation matrices for each participant's recall topic
744 proportions matrix (Fig. 2D, Supp. Fig. 5). However, to correct for length differences and poten-
745 tial non-linear transformations between viewing time and recall time, we first used dynamic time
746 warping⁶⁴ to temporally align participants' recall topic proportions matrices with the episode topic
747 proportions matrix. An example correlation matrix before and after warping is shown in Fig. 8B.
748 This yielded a 1976×1976 correlation matrix for the episode template and for each participant's
749 recall template.

750 The temporal structure of the episode's content (as described by our model) is captured in the
751 block-diagonal structure of the episode's temporal correlation matrix (e.g., Figs. 2B, 8A), with time
752 periods of thematic stability represented as dark blocks of varying sizes. Inspecting the episode
753 correlation matrix suggests that the episode's semantic content is highly temporally specific (i.e., the
754 correlations between topic vectors from distant timepoints are almost all near zero). By contrast, the
755 activity patterns of individual (cubes of) voxels can encode relatively limited information on their
756 own, and their activity frequently contributes to multiple separate functions⁶⁵⁻⁶⁸. By nature, these
757 two attributes give rise to similarities in activity across large timescales that may not necessarily
758 reflect a single task. To identify brain regions whose shifts in activity patterns mirrored shifts in the
759 semantic content of the episode or recalls, we restricted the temporal correlations we considered to
760 the timescale of semantic information captured by our model. Specifically, we isolated the upper
761 triangle of the episode correlation matrix and created a "proximal correlation mask" that included
762 only diagonals from the upper triangle of the episode correlation matrix up to the first diagonal that
763 contained no positive correlations. Applying this mask to the full episode correlation matrix was
764 equivalent to excluding diagonals beyond the corner of the largest diagonal block. In other words,

765 the timescale of temporal correlations we considered corresponded to the longest period of thematic
766 stability in the episode, and by extension the longest period of thematic stability in participants'
767 recalls and the longest period of stability we might expect to see in voxel activity arising from
768 processing or encoding episode content. Figure 8 shows this proximal correlation mask applied
769 to the temporal correlation matrices for the episode, an example participant's (warped) recall, and
770 an example cube of voxels from our searchlight analyses.

771 To determine which (cubes of) voxel responses matched the episode template, we correlated the
772 proximal diagonals from the upper triangle of the voxel correlation matrix for each cube with the
773 proximal diagonals from episode template matrix⁶⁹. This yielded, for each participant, a voxelwise
774 map of correlation values. We then performed a one-sample *t*-test on the distribution of (Fisher
775 *z*-transformed) correlations at each voxel, across participants. This resulted in a value for each
776 voxel (cube), describing how reliably its timecourse followed that of the episode.

777 We further sought to ensure that our analysis identified regions where the activations' temporal
778 structure specifically reflected that of the episode, rather than regions whose activity was simply
779 autocorrelated at a timescale similar to the episode template's diagonal. To achieve this, we used
780 a phase shift-based permutation procedure, whereby we circularly shifted the episode's topic
781 proportions matrix by a random number of timepoints (rows), computed the resulting "null"
782 episode template, and re-ran the searchlight analysis, in full. (For each of the 100 permutations, the
783 same random shift was used for all participants). We *z*-scored the observed (unshifted) result at
784 each voxel against the distribution of permutation-derived "null" results, and estimated a *p*-value
785 by computing the proportion of shifted results that yielded larger values. To create the map in
786 Figure 8C, we thresholded out any voxels whose similarity to the unshifted episode's structure fell
787 below the 95th percentile of the permutation-derived similarity results.

788 We used an analogous procedure to identify which voxels' responses reflected the recall tem-
789 plates. For each participant, we correlated the proximal diagonals from the upper triangle of the
790 correlation matrix for each cube of voxels with the proximal diagonals from the upper triangle
791 of their (time-warped) recall correlation matrix. As in the episode template analysis, this yielded
792 a voxelwise map of correlation coefficients for each participant. However, whereas the episode

793 analysis compared every participant's responses to the same template, here the recall templates
794 were unique for each participant. As in the analysis described above, we *t*-scored the (Fisher *z*-
795 transformed) voxelwise correlations, and used the same permutation procedure we developed for
796 the episode responses to ensure specificity to the recall timeseries and assign significance values.
797 To create the map in Figure 8D we again thresholded out any voxels whose scores were below the
798 95th percentile of the permutation-derived null distribution.

799 **Neurosynth decoding analyses**

800 Neurosynth³⁹ parses a massive online database of over 14000 neuroimaging studies and constructs
801 meta-analysis images for over 13000 psychology- and neuroscience-related terms, based on NIfTI
802 images accompanying studies where those terms appear at a high frequency. Given a novel image
803 (tagged with its value type; e.g., *z*-, *t*-, *F*- or *p*-statistics), Neurosynth returns a list of terms whose
804 meta-analysis images are most similar. Our permutation procedure yielded, for each of the two
805 searchlight analyses, a voxelwise map of *z*-values. These maps describe the extent to which each
806 voxel specifically reflected the temporal structure of the episode or individuals' recalls (i.e., relative
807 to the null distributions of phase-shifted values). We inputted the two statistical maps described
808 above to Neurosynth to create a list of the 10 most representative terms for each map.

809 **Data availability**

810 The fMRI data we analyzed are available online at:

811 <https://dataspace.princeton.edu/jspui/handle/88435/dsp01nz8062179>

812 The behavioral data is available at:

813 <https://github.com/ContextLab/sherlock-topic-model-paper/tree/master/data/raw>

814 **Code availability**

815 All of our analysis code may be downloaded from:

816 <https://github.com/ContextLab/sherlock-topic-model-paper>

817 **References**

- 818 [1] Murdock, B. B. The serial position effect of free recall. *Journal of Experimental Psychology* **64**,
819 482–488 (1962).
- 820 [2] Kahana, M. J. Associative retrieval processes in free recall. *Memory & Cognition* **24**, 103–109
821 (1996).
- 822 [3] Yonelinas, A. P. The nature of recollection and familiarity: A review of 30 years of research.
823 *Journal of Memory and Language* **46**, 441–517 (2002).
- 824 [4] Kahana, M. J. *Foundations of Human Memory* (Oxford University Press, New York, NY, 2012).
- 825 [5] Koriat, A. & Goldsmith, M. Memory in naturalistic and laboratory contexts: distinguish-
826 ing accuracy-oriented and quantity-oriented approaches to memory assessment. *Journal of*
827 *Experimental Psychology: General* **123**, 297–315 (1994).
- 828 [6] Huk, A., Bonnen, K. & He, B. J. Beyond trial-based paradigms: continuous behavior, ongoing
829 neural activity, and naturalistic stimuli. *Journal of Neuroscience* **10.1523/JNEUROSCI.1920-**
830 **17.2018** (2018).
- 831 [7] Lerner, Y., Honey, C. J., Silbert, L. J. & Hasson, U. Topographic mapping of a hierarchy of
832 temporal receptive windows using a narrated story. *Journal of Neuroscience* **31**, 2906–2915
833 (2011).
- 834 [8] Manning, J. R. Episodic memory: mental time travel or a quantum ‘memory wave’ function?
835 *PsyArXiv* doi:10.31234/osf.io/6zjwb (2019).
- 836 [9] Manning, J. R. Context reinstatement. In Kahana, M. J. & Wagner, A. D. (eds.) *Handbook of*
837 *Human Memory* (Oxford University Press, 2020).

- 838 [10] Howard, M. W. & Kahana, M. J. A distributed representation of temporal context. *Journal of*
839 *Mathematical Psychology* **46**, 269–299 (2002).
- 840 [11] Howard, M. W. *et al.* A unified mathematical framework for coding time, space, and sequences
841 in the medial temporal lobe. *Journal of Neuroscience* **34**, 4692–4707 (2014).
- 842 [12] Manning, J. R., Norman, K. A. & Kahana, M. J. The role of context in episodic memory. In
843 Gazzaniga, M. (ed.) *The Cognitive Neurosciences, Fifth edition*, 557–566 (MIT Press, 2015).
- 844 [13] Ranganath, C. & Ritchey, M. Two cortical systems for memory-guided behavior. *Nature*
845 *Reviews Neuroscience* **13**, 713 – 726 (2012).
- 846 [14] Zacks, J. M., Speer, N. K., Swallow, K. M., Braver, T. S. & Reynolds, J. R. Event perception: a
847 mind-brain perspective. *Psychological Bulletin* **133**, 273–293 (2007).
- 848 [15] Zwaan, R. A. & Radvansky, G. A. Situation models in language comprehension and memory.
849 *Psychological Bulletin* **123**, 162 – 185 (1998).
- 850 [16] Radvansky, G. A. & Zacks, J. M. Event boundaries in memory and cognition. *Curr Opin Behav*
851 *Sci* **17**, 133–140 (2017).
- 852 [17] Brunec, I. K., Moscovitch, M. M. & Barense, M. D. Boundaries shape cognitive representations
853 of spaces and events. *Trends in Cognitive Sciences* **22**, 637–650 (2018).
- 854 [18] Heusser, A. C., Ezzyat, Y., Shiff, I. & Davachi, L. Perceptual boundaries cause mnemonic
855 trade-offs between local boundary processing and across-trial associative binding. *Journal of*
856 *Experimental Psychology Learning, Memory, and Cognition* **44**, 1075–1090 (2018).
- 857 [19] Clewett, D. & Davachi, L. The ebb and flow of experience determines the temporal structure
858 of memory. *Curr Opin Behav Sci* **17**, 186–193 (2017).
- 859 [20] Ezzyat, Y. & Davachi, L. What constitutes an episode in episodic memory? *Psychological*
860 *Science* **22**, 243–252 (2011).

- 861 [21] DuBrow, S. & Davachi, L. The influence of contextual boundaries on memory for the sequential
862 order of events. *Journal of Experimental Psychology: General* **142**, 1277–1286 (2013).
- 863 [22] Tompary, A. & Davachi, L. Consolidation promotes the emergence of representational overlap
864 in the hippocampus and medial prefrontal cortex. *Neuron* **96**, 228–241 (2017).
- 865 [23] Chen, J. *et al.* Shared memories reveal shared structure in neural activity across individuals.
866 *Nature Neuroscience* **20**, 115 (2017).
- 867 [24] Blei, D. M., Ng, A. Y. & Jordan, M. I. Latent dirichlet allocation. *Journal of Machine Learning*
868 *Research* **3**, 993 – 1022 (2003).
- 869 [25] Rabiner, L. A tutorial on Hidden Markov Models and selected applications in speech recog-
870 nition. *Proceedings of the IEEE* **77**, 257–286 (1989).
- 871 [26] Baldassano, C. *et al.* Discovering event structure in continuous narrative perception and
872 memory. *Neuron* **95**, 709–721 (2017).
- 873 [27] Blei, D. M. & Lafferty, J. D. Dynamic topic models. In *Proceedings of the 23rd International*
874 *Conference on Machine Learning, ICML '06*, 113–120 (ACM, New York, NY, US, 2006).
- 875 [28] Manning, J. R., Polyn, S. M., Baltuch, G., Litt, B. & Kahana, M. J. Oscillatory patterns in
876 temporal lobe reveal context reinstatement during memory search. *Proceedings of the National*
877 *Academy of Sciences, USA* **108**, 12893–12897 (2011).
- 878 [29] Howard, M. W., Viskontas, I. V., Shankar, K. H. & Fried, I. Ensembles of human MTL neurons
879 “jump back in time” in response to a repeated stimulus. *Hippocampus* **22**, 1833–1847 (2012).
- 880 [30] Atkinson, R. C. & Shiffrin, R. M. Human memory: A proposed system and its control
881 processes. In Spence, K. W. & Spence, J. T. (eds.) *The psychology of learning and motivation*,
882 vol. 2, 89–105 (Academic Press, New York, 1968).
- 883 [31] Postman, L. & Phillips, L. W. Short-term temporal changes in free recall. *Quarterly Journal of*
884 *Experimental Psychology* **17**, 132–138 (1965).

- 885 [32] Welch, G. B. & Burnett, C. T. Is primacy a factor in association-formation. *American Journal of*
886 *Psychology* **35**, 396–401 (1924).
- 887 [33] Polyn, S. M., Norman, K. A. & Kahana, M. J. A context maintenance and retrieval model of
888 organizational processes in free recall. *Psychological Review* **116**, 129–156 (2009).
- 889 [34] Manning, J. R. & Kahana, M. J. Interpreting semantic clustering effects in free recall. *Memory*
890 **20**, 511–517 (2012).
- 891 [35] Heusser, A. C., Ziman, K., Owen, L. L. W. & Manning, J. R. HyperTools: a Python toolbox for
892 gaining geometric insights into high-dimensional data. *Journal of Machine Learning Research*
893 **18**, 1–6 (2018).
- 894 [36] McInnes, L., Healy, J. & Melville, J. UMAP: Uniform manifold approximation and projection
895 for dimension reduction. *arXiv* **1802** (2018).
- 896 [37] Mueller, A. *et al.* WordCloud 1.5.0: a little word cloud generator in Python. *Zenodo*
897 <https://zenodo.org/record/1322068#.W4tPKZKNKh24> (2018).
- 898 [38] Paller, K. A. & Wagner, A. D. Observing the transformation of experience into memory. *Trends*
899 *in Cognitive Sciences* **6**, 93–102 (2002).
- 900 [39] Yarkoni, T., Poldrack, R. A., Nichols, T. E., Van Essen, D. C. & Wager, T. D. Large-scale
901 automated synthesis of human functional neuroimaging data. *Nature Methods* **8**, 665 (2011).
- 902 [40] Bellmund, J. L. S., Gärdenfors, P., Moser, E. I. & Doeller, C. F. Navigating cognition: spatial
903 codes for human thinking. *Science* **362** (2018).
- 904 [41] Bellmund, J. L. S. *et al.* Deforming the metric of cognitive maps distorts memory. *Nature*
905 *Human Behavior* **4**, 177–188 (2020).
- 906 [42] Constantinescu, A. O., O'Reilly, J. X. & Behrens, T. E. J. Organizing conceptual knowledge in
907 humans with a gridlike code. *Science* **352**, 1464–1468 (2016).

- 908 [43] Gilboa, A. & Marlatte, H. Neurobiology of schemas and schema-mediated memory. *Trends*
909 *Cogn Sci* **21**, 618–631 (2017).
- 910 [44] Baldassano, C., Hasson, U. & Norman, K. A. Representation of real-world event schemas
911 during narrative perception. *Journal of Neuroscience* **38**, 9689–9699 (2018).
- 912 [45] Huth, A. G., Nisimoto, S., Vu, A. T. & Gallant, J. L. A continuous semantic space describes the
913 representation of thousands of object and action categories across the human brain. *Neuron*
914 **76**, 1210–1224 (2012).
- 915 [46] Huth, A. G., de Heer, W. A., Griffiths, T. L., Theunissen, F. E. & Gallant, J. L. Natural speech
916 reveals the semantic maps that tile human cerebral cortex. *Nature* **532**, 453–458 (2016).
- 917 [47] Gagnepain, P. *et al.* Collective memory shapes the organization of individual memories in the
918 medial prefrontal cortex. *Nature Human Behavior* **4**, 189–200 (2020).
- 919 [48] Simony, E., Honey, C. J., Chen, J. & Hasson, U. Dynamic reconfiguration of the default mode
920 network during narrative comprehension. *Nature Communications* **7**, 1–13 (2016).
- 921 [49] Zadbood, A., Chen, J., Leong, Y. C., Norman, K. A. & Hasson, U. How we transmit memories
922 to other brains: Constructing shared neural representations via communication. *Cereb Cortex*
923 **27**, 4988–5000 (2017).
- 924 [50] Simony, E. & Chang, C. Analysis of stimulus-induced brain dynamics during naturalistic
925 paradigms. *NeuroImage* **216**, 116461 (2020).
- 926 [51] Landauer, T. K. & Dumais, S. T. A solution to Plato’s problem: the latent semantic analysis
927 theory of acquisition, induction, and representation of knowledge. *Psychological Review* **104**,
928 211–240 (1997).
- 929 [52] Mikolov, T., Chen, K., Corrado, G. & Dean, J. Efficient estimation of word representations in
930 vector space. *arXiv* **1301.3781** (2013).
- 931 [53] Cer, D. *et al.* Universal sentence encoder. *arXiv* **1803.11175** (2018).

- 932 [54] Radford, A. *et al.* Language models are unsupervised multitask learners. *OpenAI Blog* **1** (2019).
- 933 [55] Brown, T. B. *et al.* Language models are few-shot learners. *arXiv* **2005.14165** (2020).
- 934 [56] Pedregosa, F. *et al.* Scikit-learn: Machine learning in Python. *Journal of Machine Learning*
935 *Research* **12**, 2825–2830 (2011).
- 936 [57] Capota, M. *et al.* Brain imaging analysis kit (2017). URL [https://doi.org/10.5281/zenodo.](https://doi.org/10.5281/zenodo.59780)
937 **59780**.
- 938 [58] Hasson, U., Yang, E., Vallines, I., Heeger, D. J. & Rubin, N. A hierarchy of temporal receptive
939 windows in human cortex. *Journal of Neuroscience* **28**, 2539–2550 (2008).
- 940 [59] Hasson, U., Chen, J. & Honey, C. J. Hierarchical process memory: memory as an integral
941 component of information processing. *Trends in Cognitive Science* **19**, 304–315 (2015).
- 942 [60] Dobrushin, R. L. Prescribing a system of random variables by conditional distributions. *Theory*
943 *of Probability & Its Applications* **15**, 458–486 (1970).
- 944 [61] Ramdas, A., Trillos, N. & Cuturi, M. On wasserstein two-sample testing and related families
945 of nonparametric tests. *Entropy* **19**, 47 (2017).
- 946 [62] Heusser, A. C., Fitzpatrick, P. C., Field, C. E., Ziman, K. & Manning, J. R. Quail: a Python
947 toolbox for analyzing and plotting free recall data. *The Journal of Open Source Software*
948 **10.21105/joss.00424** (2017).
- 949 [63] Fisher, R. A. *Statistical Methods for Research Workers* (Oliver and Boyd, 1925).
- 950 [64] Berndt, D. J. & Clifford, J. Using dynamic time warping to find patterns in time series. In
951 *KDD workshop*, vol. 10, 359–370 (1994).
- 952 [65] Freedman, D., Riesenhuber, M., Poggio, T. & Miller, E. Categorical representation of visual
953 stimuli in the primate prefrontal cortex. *Science* **291**, 312–316 (2001).
- 954 [66] Sigman, M. & Dehaene, S. Brain mechanisms of serial and parallel processing during dual-task
955 performance. *Journal of Neuroscience* **28**, 7585–7589 (2008).

- 956 [67] Charron, S. & Koechlin, E. Divided representations of current goals in the human frontal
957 lobes. *Science* **328**, 360–363 (2010).
- 958 [68] Rishel, C. A., Huang, G. & Freedman, D. J. Independent category and spatial encoding in
959 parietal cortex. *Neuron* **77**, 969–979 (2013).
- 960 [69] Kriegeskorte, N., Mur, M. & Bandettini, P. Representational similarity analysis – connecting
961 the branches of systems neuroscience. *Frontiers in Systems Neuroscience* **2**, 1 – 28 (2008).

962 **Acknowledgements**

963 We thank Luke Chang, Janice Chen, Chris Honey, Caroline Lee, Lucy Owen, Emily Whitaker,
964 Xinming Xu, and Kirsten Ziman for feedback and scientific discussions. We also thank Janice
965 Chen, Yuan Chang Leong, Chris Honey, Chung Yong, Kenneth Norman, and Uri Hasson for
966 sharing the data used in our study. Our work was supported in part by NSF EPSCoR Award
967 Number 1632738. The content is solely the responsibility of the authors and does not necessarily
968 represent the official views of our supporting organizations. The funders had no role in study
969 design, data collection and analysis, decision to publish or preparation of the manuscript.

970 **Author contributions**

971 Conceptualization: A.C.H. and J.R.M.; Methodology: A.C.H., P.C.F. and J.R.M.; Software: A.C.H.,
972 P.C.F. and J.R.M.; Analysis: A.C.H., P.C.F. and J.R.M.; Writing, Reviewing, and Editing: A.C.H.,
973 P.C.F. and J.R.M.; Supervision: J.R.M.

974 **Competing interests**

975 The authors declare no competing interests.

Figure 1: Topic weights in episode and recall content. We used detailed, hand-generated annotations describing each manually identified time segment from the episode to fit a topic model. Three example frames from the episode (first row) are displayed, along with their descriptions from the corresponding episode annotation (second row) and an example participant’s recall transcript (third row). We used the topic model (fit to the episode annotations) to estimate topic vectors for each moment of the episode and each sentence of participants’ recalls. Example topic vectors are displayed in the bottom row (blue: episode annotations; green: example participant’s recalls). Three topic dimensions are shown (the highest-weighted topics for each of the three example scenes, respectively), along with the 10 highest-weighted words for each topic. Supplementary Figure 2 provides a full list of the top 10 words from each of the discovered topics.

Figure 2: Modeling naturalistic stimuli and recalls. All panels: darker colors indicate greater values; range: [0, 1]. **A.** Topic vectors ($K = 100$) for each of the 1976 episode timepoints. **B.** Timepoint-by-timepoint correlation matrix of the topic vectors displayed in Panel A. Event boundaries discovered by the HMM are denoted in yellow (30 events detected). **C.** Average topic vectors for each of the 30 episode events. **D.** Topic vectors for each of 265 sliding windows of sentences spoken by an example participant while recalling the episode. **E.** Timepoint-by-timepoint correlation matrix of the topic vectors displayed in Panel D. Event boundaries detected by the HMM are denoted in yellow (22 events detected). For similar plots for all participants, see Supplementary Figure 5. **F.** Average topic vectors for each of the 22 recall events from the example participant. **G.** Correlations between the topic vectors for every pair of episode events (Panel C) and recall events (from the example participant; Panel F). For similar plots for all participants, see Supplementary Figure 6. **H.** Average correlations between each pair of episode events and recall events (across all 17 participants). To create the figure, each recalled event was assigned to the episode event with the most correlated topic vector (yellow boxes in panels G and H).

Figure 3: Naturalistic extensions of classic list-learning memory analyses. **A.** The probability of first recall as a function of the serial position of the event in the episode. **B.** The probability of recalling each event, conditioned on having most recently recalled the event *lag* events away in the episode. **C.** The proportion of participants who recalled each event, as a function of the serial position of the events in the episode. All panels: error ribbons denote the bootstrap-estimated 95% confidence interval.

Figure 4: Novel content-based metrics of naturalistic memory: precision and distinctiveness. **A.** The episode-recall correlation matrix for an example participant (P17), chosen for their large number of recall events (for analogous figures for other participants, see Supp. Fig. 5). The yellow boxes highlight the maximum correlation in each column. The example participant’s overall precision score was computed as the average across the (Fisher z-transformed) correlation values in the yellow boxes. Their distinctiveness score was computed as the average (over recall events) of the z-scored (within column) event precisions. **B.** The across-participants (Pearson’s) correlation between precision and hand-counted number of recalled scenes. **C.** The correlation between distinctiveness and hand-counted number of recalled scenes. **D.** The correlation between precision and the number of recalled episode events, as determined by our model. **E.** The correlation between distinctiveness and the number of recalled episode events, as determined by our model.

Figure 5: Precision reflects the completeness of recall, whereas distinctiveness reflects recall specificity. A. Recall precision by episode event. Grey violin plots display kernel density estimates for the distribution of recall precision scores for a single episode event. Colored dots within each violin plot represent individual participants' recall precisions for the given event. **B.** Recall distinctiveness by episode event, analogous to Panel A. **C.** The set of "Narrative Details" episode annotations²³ comprising an example episode event (22) identified by the HMM. Each action or feature is highlighted in a different color. **D.** Sentences comprising the most precise (P17) and least precise (P6) participants' recalls of episode event 21. Descriptions of specific actions or features reflecting those highlighted in Panel B are highlighted in the corresponding color. The text highlighted in gray denotes a (rare) false recall. The unhighlighted text denotes correctly recalled information about other episode events. **E.** The sets of "Narrative Details" episode annotations²³ for scenes comprising episode events described by the example participants in Panel F. Each event's text is highlighted in a different color. **F.** The sentences comprising the most distinctive (P9) and least distinctive (P6) participants' recalls of episode event 21. Sections of recall describing each episode event in Panel E are highlighted with the corresponding color.

Figure 6: Trajectories through topic space capture the dynamic content of the episode and recalls. All panels: the topic proportion matrices have been projected onto a shared two-dimensional space using UMAP. **A.** The two-dimensional topic trajectory taken by the episode of *Sherlock*. Each dot indicates an event identified using the HMM (see *Methods*); the dot colors denote the order of the events (early events are in red; later events are in blue), and the connecting lines indicate the transitions between successive events. **B.** The average two-dimensional trajectory captured by participants' recall sequences, with the same format and coloring as the trajectory in Panel A. To compute the event positions, we matched each recalled event with an event from the original episode (see *Results*), and then we averaged the positions of all events with the same label. The arrows reflect the average transition direction through topic space taken by any participants whose trajectories crossed that part of topic space; blue denotes reliable agreement across participants via a Rayleigh test ($p < 0.05$, corrected). For additional detail see *Methods* and Supplementary Figure 4. **C.** The recall topic trajectories (gray) taken by each individual participant (P1–P17). The episode's trajectory is shown in black for reference. Here, events (dots) are colored by their matched episode event (Panel A).

Figure 7: Language used in the most and least precisely remembered events. **A.** Average precision (episode event-recall event topic vector correlation) across participants for each episode event. Here we defined each episode event's precision for each participant as the correlation between its topic vector and the most-correlated recall event's topic vector from that participant. Error bars denote bootstrap-derived across-participant 95% confidence intervals. The stars denote the three most precisely remembered events (green) and least precisely remembered events (red). **B.** Wordles comprising the top 200 highest-weighted words reflected in the weighted-average topic vector across episode events. Green: episode events were weighted by their precision (Panel A). Red: episode events were weighted by the inverse of their precision. **C.** The set of all episode and recall events is projected onto the two-dimensional space derived in Figure 6. The dots outlined in black denote episode events (dot size is proportional to each event's average precision). The dots without black outlines denote individual recall events from each participant. All dots are colored using the same scheme as Figure 6A. Wordles for several example events are displayed (green: three most precisely remembered events; red: three least precisely remembered events). Within each circular wordle, the left side displays words associated with the topic vector for the episode event, and the right side displays words associated with the (average) recall event topic vector, across all recall events matched to the given episode event.

Figure 8: Brain structures that underlie the transformation of experience into memory. **A.** We isolated the proximal diagonals from the upper triangle of the episode correlation matrix, and applied this same diagonal mask to the voxel response correlation matrix for each cube of voxels in the brain. We then searched for brain regions whose activation timeseries consistently exhibited a similar proximal correlational structure to the episode model, across participants. **B.** We used dynamic time warping⁶⁴ to align each participant's recall timeseries to the TR timeseries of the episode. We then computed the temporal correlation matrix of each participant's warped recalls. Next, we applied the same diagonal mask used in Panel A to isolate the proximal temporal correlations and searched for brain regions whose activation timeseries for each participant consistently exhibited a similar proximal correlational structure to that participant's recalls. **C.** We identified a network of regions sensitive to the narrative structure of participants' ongoing experience. The map shown is thresholded at $p < 0.05$, corrected. The top ten Neurosynth terms displayed in the panel were computed using the unthresholded map. **D.** We also identified a network of regions sensitive to how individuals would later structure the episode's content in their recalls. The map shown is thresholded at $p < 0.05$, corrected. The top ten Neurosynth terms displayed in the panel were computed using the unthresholded map.